

Building Resilience into Go Services

By Max Collier

Topics for today

A recounting of production service failures 

Consider error handling and vulnerable in service design

Open source solution to address concerns

Today's Service

Alert Sender...

Example: *Vessel has entered area*

Vessel <Endeavour> Has entered Area <New Zealand
400mn> At <01/.01/.2025 00:00 NZST>

Example: *Vessel has exited area*

Vessel <Endeavour> Has exited Area <New Zealand
400mn> At <02/.01/.2025 00:00 NZST>

Example: Vessel missing AIS in area

Vessel <Endeavour> has stopped reporting AIS within
Area <New Zealand 400mn> From <02/.01/.2025 00:00
NZST> for over <X duration>

Architecture

How do we represent this in a system?

Alerting Framework

What do I need?

- Alert Rules
- Vessel Identities
- Users with timezones
- Vessel Events (Activity)
- Additional Vessel Information

Alerting Framework

Alert Sender

Alert Sender

Alert sender complexity

Architecture constraints

- Long running processes: one deployment per alert-types (i.e digest, or realtime)
- Core process loop: handles all alert-notifications across organisations / users.
- Alerts Framework: requires notifications to be queued to send before marking as complete.

So, in other words

Alert sender has no redundancy

...

Game Time!

Sherlock...

Ready?

First Outage

No error handling

What went wrong?

Alert Sender

```
func (s *service) processAlerts(notifications []*starboard.Notifications) {
    // ...

    users, err := starboard.UsersFromIDs(ctx, userIDs)
    If err != nil {
        // TODO: log error
        return
    }

    vessels, err := starboard.VesselsFromVesselIDs(ctx, vesselIDs)
    If err != nil {
        // TODO: log error
        return
    }

    // ...
}
```

Alert Sender

```
func (s *service) ProcessAlerts(notifications []*starboard.Notifications) error {
    // ...

    users, err := starboard.UsersFromIDs(ctx, userIDs)
    If err != nil {
        log.Error(err).Msg("query_failed_user_profiles")

        return err
    }

    vessels, err := starboard.VesselsFromVesselIDs(ctx, vesselIDs)
    If err != nil {
        log.Error(err).Msg("query_failed_vessel_profiles")

        return err
    }

    // ...
}
```

Consequences

- Customers let us know they weren't receiving emails
- Errors had been occurring in our system but we had no visibility into when or where they were occurring

Postmortem

- Log errors :faceplam:
- Fix issue that caused the problem
- Apologies to customers, explain, fix, move on.

Next Up

Second Outage

Panic occurrences

Considerations

- Panics are not protected around go routine boundaries
- Long running service requires it's own panic handling
- HTTP / gRPC services provide default panic handlers
- Panic handling can be considered an afterthought from initial service design

Alert Sender V1

```
func (s *service) run() {  
    // setup ...  
  
    for {  
        select {  
            case <-shutdown:  
                os.Exit(0)  
            case <-s.ticker.C:  
  
                notifications, err := s.PendingAlerts(ctx)  
                if err != nil {  
                    // TODO: logging  
                    continue  
                }  
  
                if len(notifications) > 0 {  
                    go s.ProcessAlerts(notifications)  
                }  
        }  
    }  
}
```

With Recovery

```
func (s *service) run() {  
    // setup ...  
  
    for {  
        select {  
            case <-shutdown:  
                os.Exit(0)  
            case <-s.ticker.C:  
  
                notifications, err := s.PendingAlerts(ctx)  
                if err != nil {  
                    log.Error(err).Msg("failed_query_pending_alerts")  
                    continue  
                }  
  
                // refactor ...  
                go func() {  
                    err := s.ProcessAlerts(notifications)  
                    if err != nil {  
                        log.Error(err).Msg("failed_process_alerts_invoke")  
                        continue  
                    }  
                    log.Error(err).Msg("successful_process_alerts_invoke")  
                }()  
        }  
    }  
}
```

With Recovery

```
// refactor ...
go func() {
    defer () {
        if r := recover(); r != nil {
            log.WithField("panic", r).Msg("panic_occured_process_alerts_invoke")
        }
    }()
}

err := s.ProcessAlerts(notifications)
if err != nil {
    log.Error(err).Msg("failed_process_alerts_invoke")

    continue
}

log.Error(err).Msg("successful_process_alerts_invoke")
}()
```

What is the language's philosophy of Go's panic behaviour?

Go Philosophy on Panic Handling

- Errors are normal; panics are exceptional.
- Go encourages explicit error handling of `error`s.
- Panics should be reserved for unexpected, or unrecoverable situations.
- Panics are not for flow control.

Types of panic occurrences

Divide by zero

Invalid type assertion

Concurrent map writes

Nil pointer dereference

Array or slice index out of range

Explicit Intent

Considerations with Go Routines?

Http handlers

```
http.HandleFunc("/", func(w http.ResponseWriter, r *http.Request) {  
    panic("oops") // Won't crash the server  
})
```

In our case

An unused fields is suddenly used

...

1. UserDB: user.isDeleted = TRUE
2. user has active alert rules/subscriptions
3. alert sender cannot query user
4. Lookup on user causes nil pointer dereference

Consequences

- Customers let us know they weren't receiving emails
- Panics brought down system for all users when emails wouldn't send for a particular user

Postmortem

- Fixed bug for deleted users
- Introduce panic recovery for handler
- Apologies to customers, explain, fix, move on.

Next Up

Third Outage

No alerts set for application errors

Story Time

Service Diagram

Consequences

- Customers let us know they weren't receiving emails
- Errors in templating alerts should fail independently
- Slack alerts for service error logs and OOM events

Postmortem

- Fixed bug with templating alerts
- Deleted pending notifications in prod older than 48hrs
- Apologies to customers, explain, fix, move on

Open Source


```
package main

import (
    "log"
    "time"

    exitmanager "github.com/mcwalrus/exit-manager"
)

func main() {
    em := exitmanager.Global()

    // Start workers
    for i := 0; i < 3; i++ {
        go criticalWorker(em, i)
    }

    // Wait for shutdown signal
    <-em.Notify()

    // Avoid exit on main routine
    select {}
}
```

```
func criticalWorker(em *exitmanager.ExitManager, id int) {
    for {
        // Acquire lock to prevent shutdown
        if err := em.AcquireShutdownLock(); err != nil {
            return
        }

        // Simulate work
        time.Sleep(3 * time.Second)
        em.ReleaseShutdownLock()

        // Check for shutdown signal
        select {
        case <-em.Notify():
            return
        case <-time.After(1 * time.Second):
        }
    }
}
```


Sentinel's Philosophy

- Treat panics as an errors
- Metrics can paint a picture of services
- Proactively recover panics for failiures in services

<https://github.com/mcwalrus/go-sentinel>

To Recap

- Handle errors
- Consider panic handlers
- Alert on service errors / failure events

Thanks for listening

